

Shawn Houser

STA 574 - Statistical Software Packages

Midterm Graduate Presentation - Using Sqlite3 and SQL databases

2019-11-05

What is a Database Management System

A tool that provides a layer of abstraction
over the data that you interact with.

Without a Database

You have to interact with a bunch of files
Imagine doing analysis on MainStreet
There will be a bunch of different data types
There could be many csv's, json, xml, or worse
You may have to deal with custom datatypes
USM may have a custom format that you need
to extract data from.

A database will solve these issues

Sqlite3 is a specific database manager (it manages the data)

Why Sqlite3

It saves all of the data in one file
You interact with this data using SQL
It can be easier to ask for data in SQL
than it would be to write code in R
It has been tried and tested
Universal, works on nearly anything
Used everywhere
Web browsers, cars, planes
Databases it produces are a single file that
can be transferred without fear of different
computers doing nasty things to the data

What does a database look like?

Collection of tables
Each table is like its own csv file
Relations between tables
Different tables can share values

Example Database

words.csv

---->

word_string	word_id
Strawberry	1
Apple	2
...	...
Saturday	93,784

counts.csv

---->

word_id	word_count
1	3992
2	42
93,784	112

Installing sqlite3

```
install.packages("RSQLite")
```

This should be all you need.

Usure if this will install sqlite3 itself,

or only allow connection if it is already installed

Some of the examples in this document may not make sense on their own. They all require datafiles to be run correctly. This file was meant to be mostly a compilation of my notes for this presentation. The presentation is meant to be interactive, and many of the results would have been shown in realtime.

The data can be found here

<https://drive.google.com/open?id=1IGfAOdZQ9DyTz8RLnj-ucnIOHG6eILYF>

You must be signed into a USM email address.

This is a small demo to show how to interact with a sqlite3 database from R, and its counterpart of interacting with a csv file.

Regular.r

```
paste("Starting at ", Sys.time())

words <-
read.csv("/Users/shawn/Desktop/ssp_midterm/data/words.
csv")
head(words)

paste("Ending at ", Sys.time())
```

Sql.r

```
paste("Starting at ", Sys.time())

# https://db.rstudio.com/databases/sqlite/
library(DBI)
database <- dbConnect(RSQLite::SQLite(),
"/Users/shawn/Desktop/ssp_midterm/data/word_database.d
b")
res <- dbSendQuery(database, "SELECT * FROM words
ORDER BY word_id asc LIMIT 6")
dbFetch(res)

paste("Ending at ", Sys.time())
```

After running both these files, we can see that the sqlite3 version ran much faster. The reason is because we don't need to read the entire file, only the parts that we want to.

Not all interactions with sqlite3 need to be stored in a file, information can be stored in the RAM of your computer and will exist only as long as the R program exists. This is useful because you will not have to keep extra files that are useless to you after computation has been completed.

Regular.r

```
subset(mtcars, cyl == 6)
```

Sql.r

```
library(DBI)

# Connects to a database that isn't on the harddrive
# Useful for something that you do not want to deal
# with afterwards, or start with a clean slate each
# time
database <- dbConnect(RSQLite::SQLite(), ":memory:")

# mtcars is the cars dataset. It is already loaded
head(mtcars)
cat("\n")

# Stores the cars dataframe into a new table called
# mtcars
dbWriteTable(database, "mtcars", mtcars)

# Shows what the table looks like
dbListFields(database, "mtcars")
cat("\n")

res <- dbSendQuery(database, "SELECT * FROM mtcars
WHERE cyl = 6")
dbFetch(res)
```

We can see that the sqlite3 version took a lot more setup than its pure R counterpart. This may be inefficient for small projects and tables like this, but for bigger problems the sqlite3 solution will be simpler.

There are functions that are very similar to R in sql. An example of this is the max function. It will iterate over all rows in a table and select the row with the maximum of that column.

Regular.r

```
paste("Starting at ", Sys.time())

counts <-
read.csv("/Users/shawn/Desktop/ssp_midterm/data/counts.csv")
max_word <- counts[which.max(counts$word_count), ]
paste("The max word_id was", max_word$word_id, "with a count
of", max_word$word_count)

paste("Ending at ", Sys.time())
```

Sql.r

```
paste("Starting at ", Sys.time())

# https://db.rstudio.com/databases/sqlite/
library(DBI)
database <- dbConnect(RSQLite::SQLite(),
"/Users/shawn/Desktop/ssp_midterm/data/word_database.d
b")
res <- dbSendQuery(database, paste(
  "SELECT word_id, max(word_count)",
  "FROM counts"
))
max_word <- dbFetch(res)

# Doesn't work because `max(word_count)` changes the
name of
# the row, you must either change the name here or in
the query
paste("The max word_id was", max_word$word_id, "with a
count of", max_word$word_count)
cat("\n")
paste("The max word_id was", max_word$word_id, "with a
count of", max_word$max)

paste("Ending at ", Sys.time())
```

Similar to 03_max, but with count. It counts the number of rows.

Regular.r

```
nrow(subset(mtcars, cyl == 6))
```

Sql.r

```
library(DBI)

database <- dbConnect(RSQLite::SQLite(), ":memory:")
dbWriteTable(database, "mtcars", mtcars)
res <- dbSendQuery(database, paste(
  "SELECT count(*)",
  "FROM mtcars",
  "WHERE cyl = 6"
))
dbFetch(res)
```

The real utility in sql is being able to work with multiple tables at once. This can become tiresome if you are working with just csv files.

Regular.r

```
paste("Starting at ", Sys.time())

counts <-
read.csv("/Users/shawn/Desktop/ssp_midterm/data/counts
.csv")
words <-
read.csv("/Users/shawn/Desktop/ssp_midterm/data/words.
.csv")

#
https://www.rdocumentation.org/packages/base/versions/
3.6.1/topics/merge
#
https://stackoverflow.com/questions/1299871/how-to-join-merge-data-frames-inner-outer-left-right
words_and_counts <- merge(counts, words, by =
"word_id")
max_word <-
words_and_counts[which.max(words_and_counts$word_count
), ]
paste(
  "The max word was `", max_word$word, "`",
  "with word_id", max_word$word_id,
  "and a count of", max_word$word_count
)

paste("Ending at ", Sys.time())
```

Sql code on next page

Sql.r

```
paste("Starting at ", Sys.time())

# https://db.rstudio.com/databases/sqlite/
library(DBI)
database <- dbConnect(RSQLite::SQLite(),
"/Users/shawn/Desktop/ssp_midterm/data/word_database.d
b")
res <- dbSendQuery(database, paste(
  "SELECT max(word_count), word_id, word",
  "FROM words",
  "INNER JOIN counts USING (word_id)"
))
max_word <- dbFetch(res)
paste(
  "The max word was `", max_word$word, "`",
  "with word_id", max_word$word_id,
  "and a count of", max_word$max
)

paste("Ending at ", Sys.time())
```

We can see that again sqlite3 is much faster than reading a csv file and doing work on that. The reason for this is the database is smart about how it stores the data. It can do operations much faster than we could.

Many times, the queries that you want to run could be very complex. When the query becomes that complex it is useful to move it into its own file. R can then read the file and use the sql in there to run a query.

Sql.r

```
library(DBI)
library(readr)

database <- dbConnect(RSQLite::SQLite(),
"/Users/shawn/Desktop/ssp_midterm/data/senators.db")
res <- dbSendQuery(database,
read_file("/Users/shawn/Desktop/ssp_midterm/06_scripts/probl
em2.sql"))
dbFetch(res)
```

Problem2.sql

```
select distinct
    stname,
    cname,
    sname,
    age,
    totrv,
    population
from senators
inner join corporations using (stname)
inner join states using (stname)
where
    totrv > (3*population)/2
    and gender = 'Female'
    and party = 'Democrat'
order by stname, cname, sname
```

Correct_output.txt

stname	cname	sname
age	totrev	population
CA	Acme Mountain Gear	Feinstein
46	50000000	29760021
CA	Rockemwell Aerospace	Feinstein
46	2147483647	29760021
IL	Da Bulls Novelty Co.	Moseley-Braun
47	24235683	11430602

```
MD      | Northeast Burro Transport | Mikulski      |
49 |      8347902 |      4781468
(4 rows)
```

As we can see we can make the whole project a lot cleaner. At this point you do not need to edit the R code to get different results from the query, and then the same query could be used in another project in possibly a different language entirely.